

# Supervised Machine Learning

## Part One: Regression



# Agenda

- Overview of Supervised Learning
- Regression Models (Algorithms)
- Model Evaluation
- Hands-on Lab

# Supervised Learning

# Definition

Supervised learning is the machine learning task of inferring a function from labeled training data. The training data consist of a set of training examples. In supervised learning, each example is a pair consisting of an input object (typically a vector) and a desired output value (also called the supervisory signal). A supervised learning algorithm analyzes the training data and produces an inferred function, which can be used for mapping new examples.

[https://en.wikipedia.org/wiki/Supervised\\_learning](https://en.wikipedia.org/wiki/Supervised_learning)

# Important Points

1. Labeled training data
2. Desired output
3. Produces an inferred function
4. Used for novel examples

# Approaches

1. Classification

2. Regression

# Regression Models

# Regression Models

- Supervised learning algorithms that estimate the relationship among variables.
- Focus is on the relationship between a dependent variable (target) and 1(+) independent variables (predictor)
- Does the dependent variable change when the independent variable(s) change?
- Common algorithms
  - Generalized linear models



# Generalized Linear Models

# Linear Models

**Linear Regression** fits a linear model to the data by adjusting a set of coefficients  $w$  to minimize the residual sum of squares between observed responses & prediction.

(1) Linear model

$$y = X\beta + \epsilon$$

(2) Objective function

$$\min_w \sum (Xw - y)^2$$

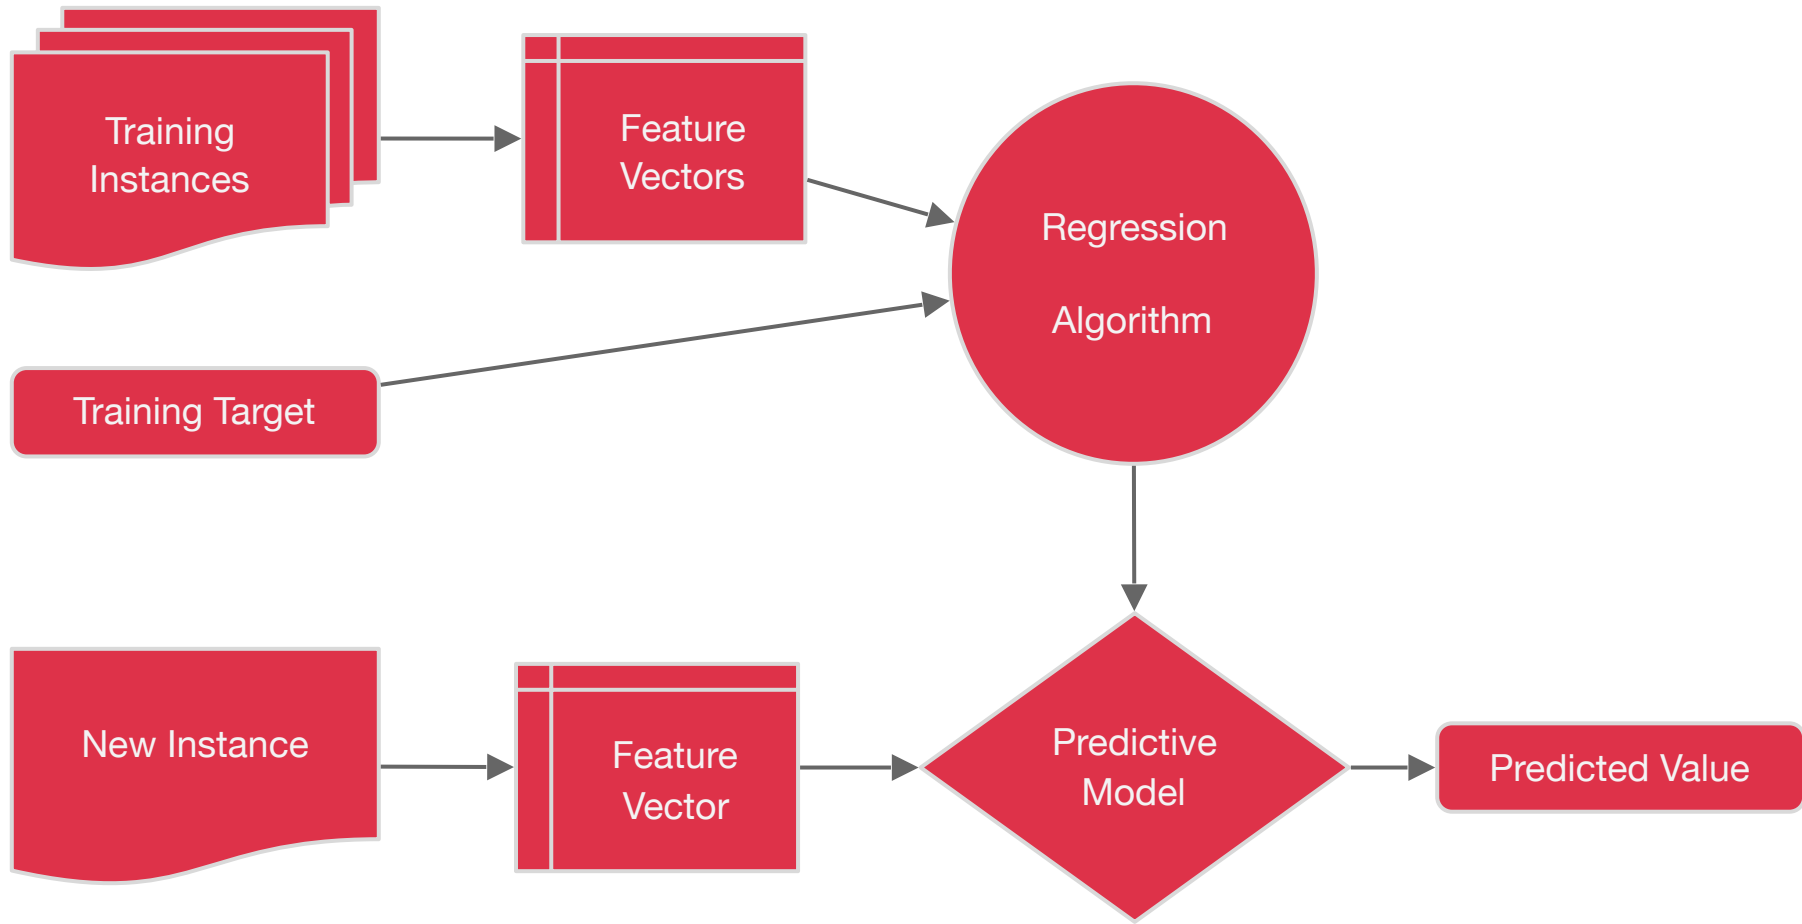
(3) Predictive model

$$\hat{y}(w, x) = w_0 + w_1x_1 + \dots + w_px_p$$

Notation:

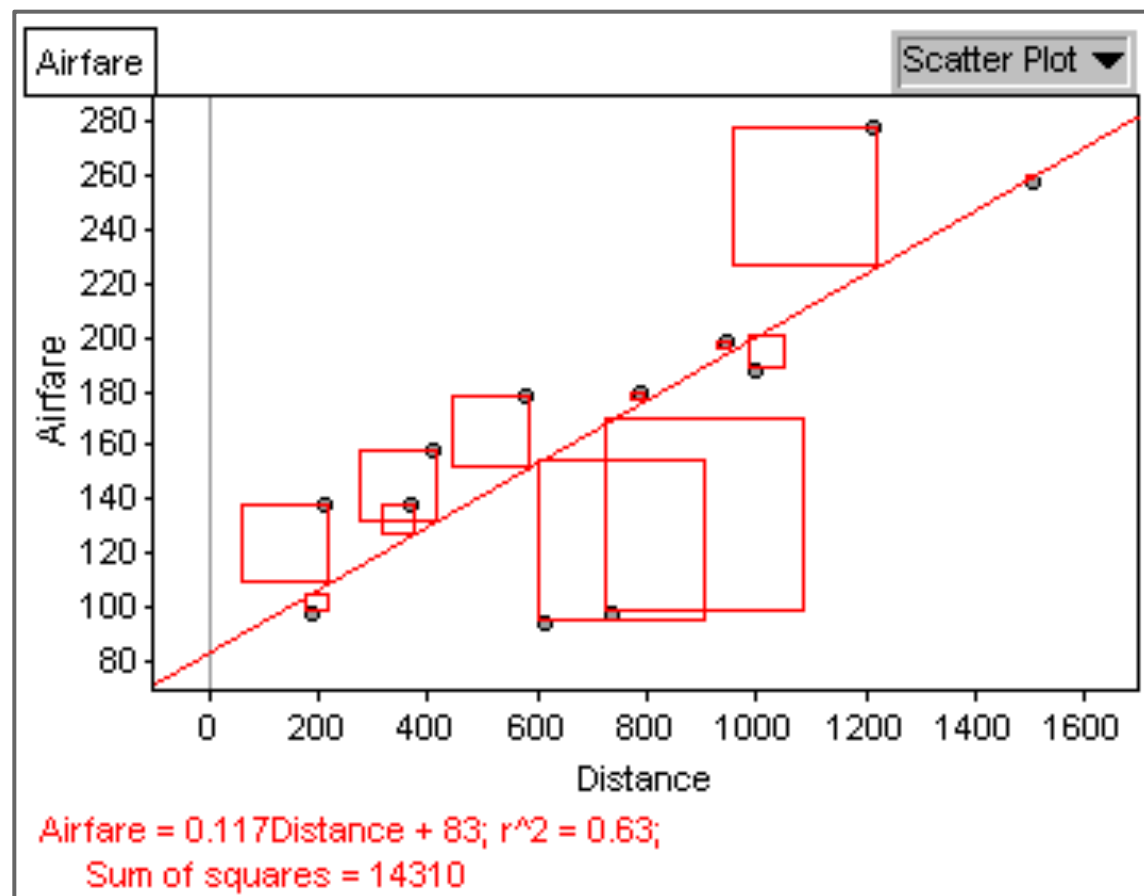
- $y$  is the observed value
- $x$  is the input variables
- $\beta$  is the set of coefficients
- $\epsilon$  is noise or randomness in observation
- $w$  is the set of weights
- $w_0$  is the ability to adjust the plane in space
- $\hat{y}$  is the predicted value

# Regression Pipeline



# Ordinary Least Squares

- Method for estimating unknown parameters in a linear regression model
- Keep adjusting parameters until minimum squared residuals (e.g. minimize some cost function).
- Relies on the independence of the model terms
- *multicollinearity*: two or more predictor variables in a multiple regression model are highly correlated, one can be linearly predicted from the others
- If this happens, the estimate becomes sensitive to error.



# Simple Regression with OLS

```
from sklearn import linear_model
from sklearn.metrics import mean_squared_error, r2_score

regr = linear_model.LinearRegression()
regr.fit(X_train, y_train)
LinearRegression(copy_X=True, fit_intercept=True, normalize=False)

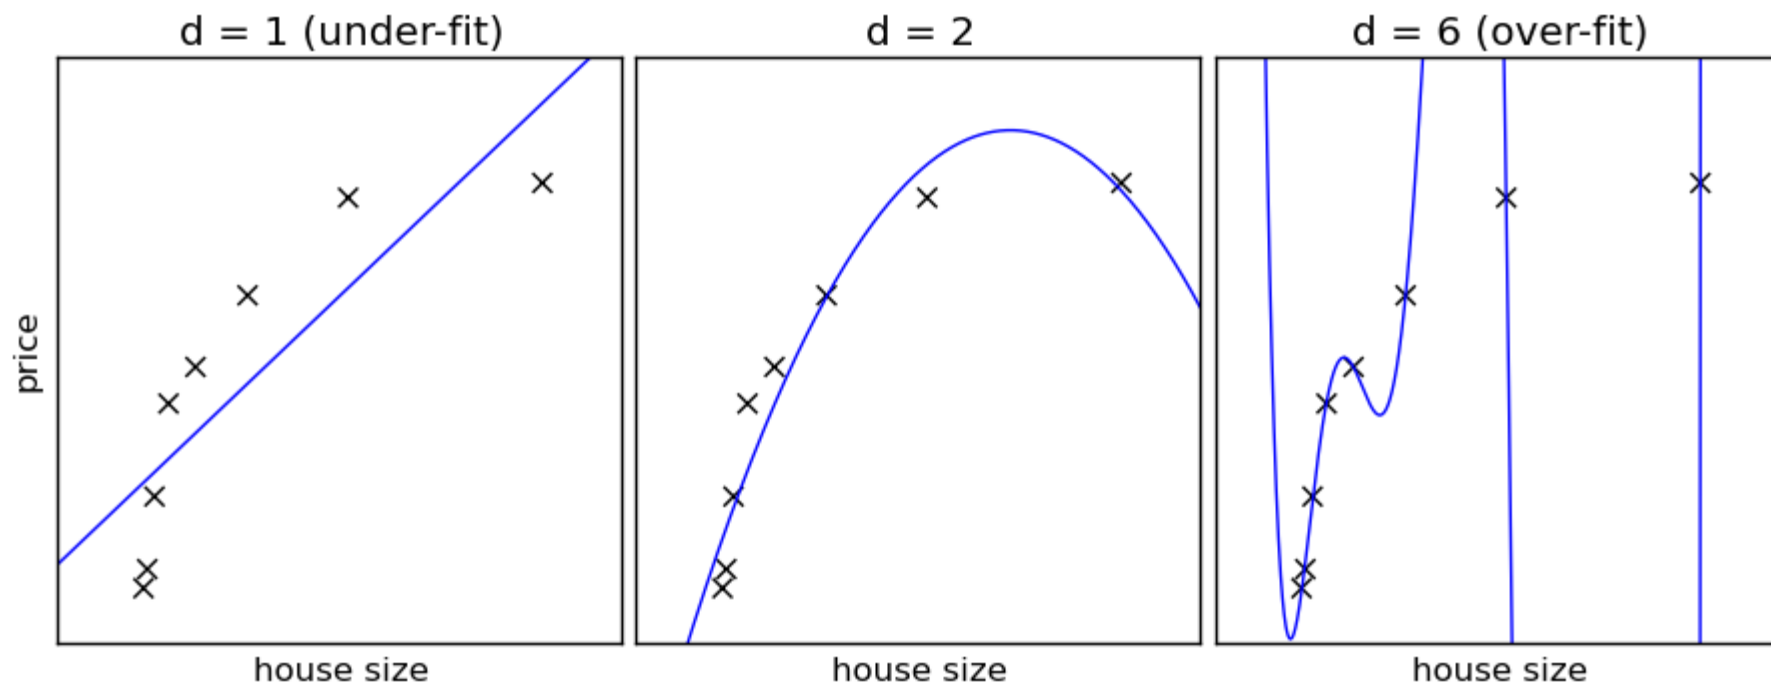
print regr.coef_
[ -6.02985639e+01  -3.02367158e+11   3.02367158e+11   6.04734316e+11
   4.17860883e+00  -3.41060763e-02   2.03234971e+01   2.15758256e-01]

print regr.intercept_
76.9490920195

print mean_squared_error(y_test, regr.predict(X_test))
7.92744075579

regr.score(X_test, y_test)  # r2_score(y_test, regr.predict(X_test))
0.92521397739317868
```

# What Can Go Wrong With Simple Linear Models?



# Regularization

- As we increase the complexity of the model we reduce the bias but increase the variance of the model.
- Variance: the tendency for the model to fit to noise (randomness) -- overfit.
- Introduce a parameter to penalize complexity in the function being minimized.



# Vector Norm

- Describes the length of the vector.
- L1: sum of the absolute values of components
- L2: euclidian distance from the origin
- $L_\infty$ : maximal absolute value component

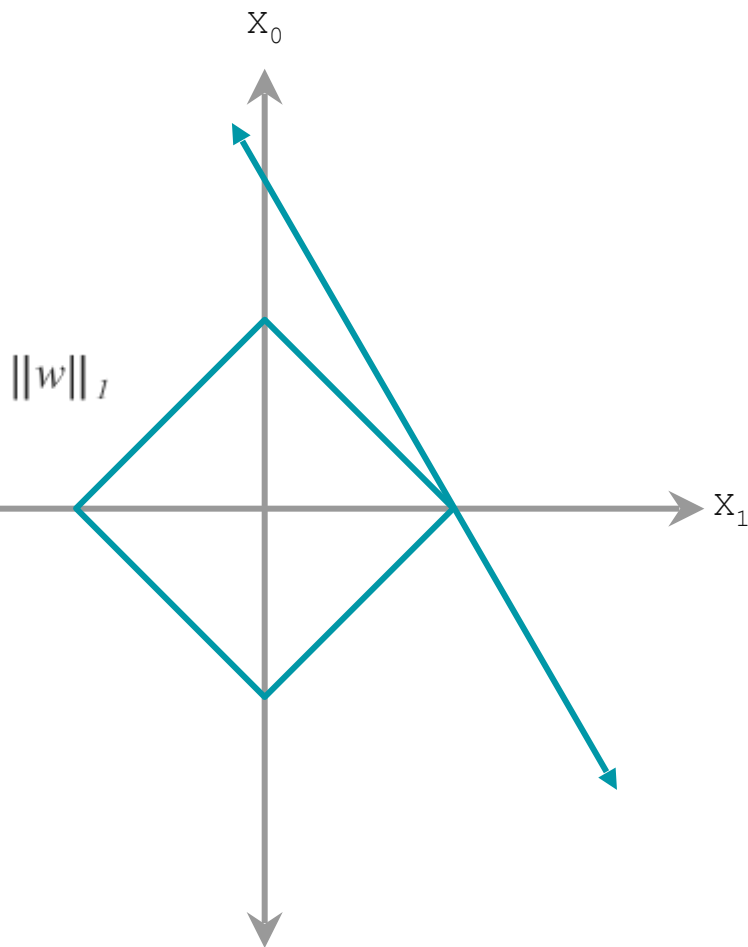
# Vector Norm

```
import numpy as np
import numpy.linalg as la

vec = np.array([-10, 3, 3, -5, -3, -2, 1, 9, 3, 4, 6,
                -8])
l1 = la.norm(vec, 1)
# 57.0

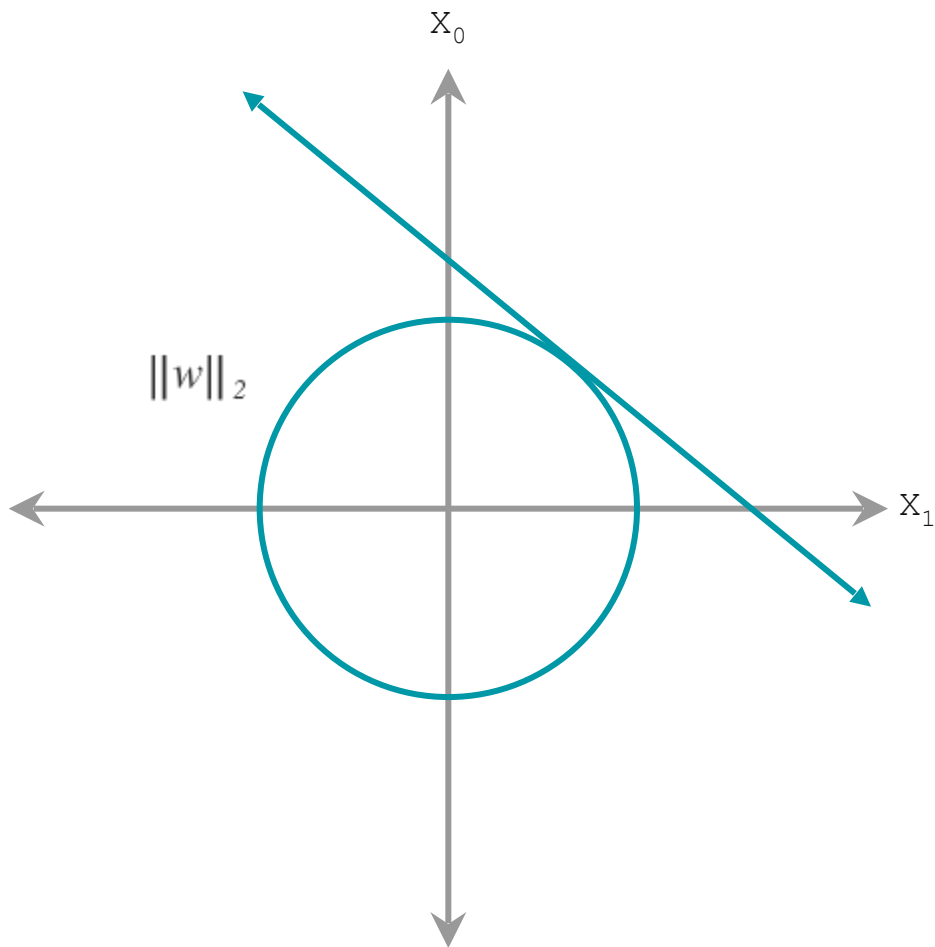
l2 = la.norm(vec, 2)
# 19.0525588833

lin = la.norm(vec, 'inf')
# 10
```



**L1 Normalization**

Possibility that a feature is eliminated by setting its coefficient equal to zero.



**L2 Normalization**

Features are kept balanced by minimizing the relative change of coefficients during learning.

# Ridge Regression

- Prevent overfit/collinearity by penalizing the size of coefficients - minimize the penalized residual sum of squares:
- Said another way, shrink the coefficients to zero.

$$\min_w \sum (Xw - y)^2 + \alpha \sum w^2$$

- Where  $\alpha > 0$  is complexity parameter that controls shrinkage. The larger  $\alpha$ , the more robust the model to collinearity.
- Alpha influences the bias/variance tradeoff: the larger the ridge alpha, the higher the bias and the lower the variance.

# Ridge Regression

```
clf = linear_model.Ridge(alpha=0.5)
clf.fit(X_train, y_train)
```

```
Ridge(alpha=0.5, copy_X=True, fit_intercept=True,
      max_iter=None, normalize=False,
      solver='auto', tol=0.001)
```

```
print mean_squared_error(y_test, clf.predict(X_test))
8.34260312032
```

```
clf.score(X_test, y_test)
0.92129741176557278
```

# Choosing alpha

We can search for the best parameter using the RidgeCV which is a form of Grid Search, but uses a more efficient form of leave-one-out cross-validation.

```
import numpy as np
n_alphas = 200
alphas = np.logspace(-10, -2, n_alphas)
clf = linear_model.RidgeCV(alphas=alphas)
clf.fit(X_train, y_train)
```

```
print clf.alpha_
0.0010843659686896108
```

```
clf.score(X_test, y_test)
0.92542477512171173
```

# Model Evaluation

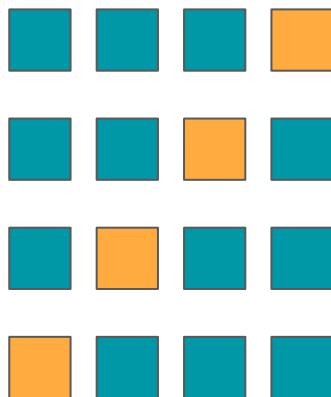
# Cross-Validation and Evaluation

- In regressions we can determine how well the model fits by computing the mean square error and the coefficient of determination.
- $MSE = np.mean((predicted - expected)^2)$
- $R^2$  is a predictor of “goodness of fit” and is a value  $\in [0,1]$  where 1 is perfect fit.



# Cross-Validation and Evaluation

In order to prevent overfit and be assured of generalizability, cross-validation fits the model on a portion of the data set and evaluates it on an unseen portion of the data set. Shuffle data, split into a large train set and smaller test set. This can be done  $K=12$  times, and scores averaged.



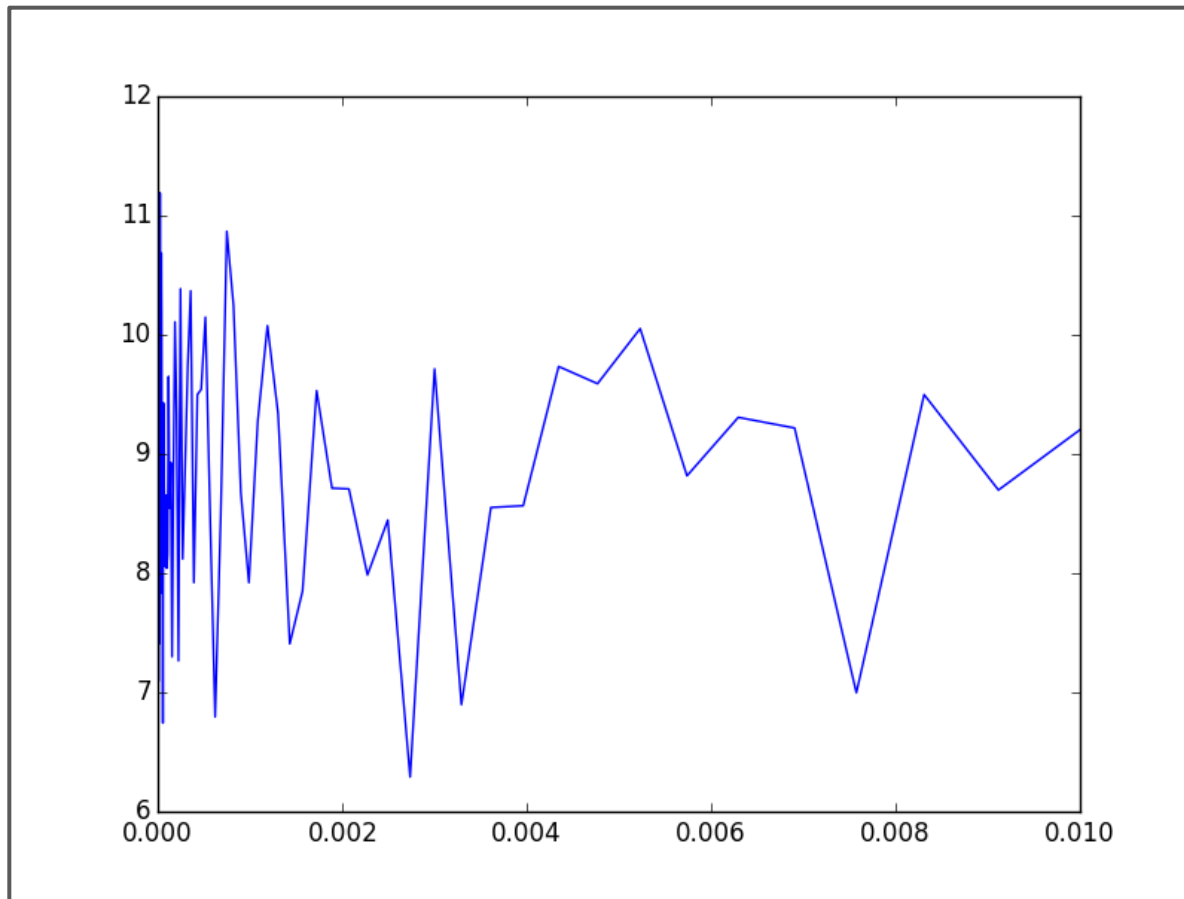
# Error As a Function of Alpha

```
clf = linear_model.Ridge(fit_intercept=False)
errors = []

for alpha in alphas:
    splits = tts(dataset.data, dataset.target('Y1'), test_size=0.2)
    X_train, X_test, y_train, y_test = splits
    clf.set_params(alpha=alpha)
    clf.fit(X_train, y_train)
    error = mean_squared_error(y_test, clf.predict(X_test))
    errors.append(error)

axe = plt.gca()
axe.plot(alphas, errors)
plt.show()
```

# Resulting Plot



# How to pick the right parameters?



# Search/Tuning

## Search Requires

- Estimator
- Parameter Space
- Method for sampling
- Cross validation scheme
- A score function

## Search Types

- Exhaustive
- Randomized
- Parallel
- Leave One Out
- Model Specific

# Creating an Exhaustive Grid Search for a Classifier

```
from sklearn.svm import SVC
from sklearn.grid_search import GridSearchCV

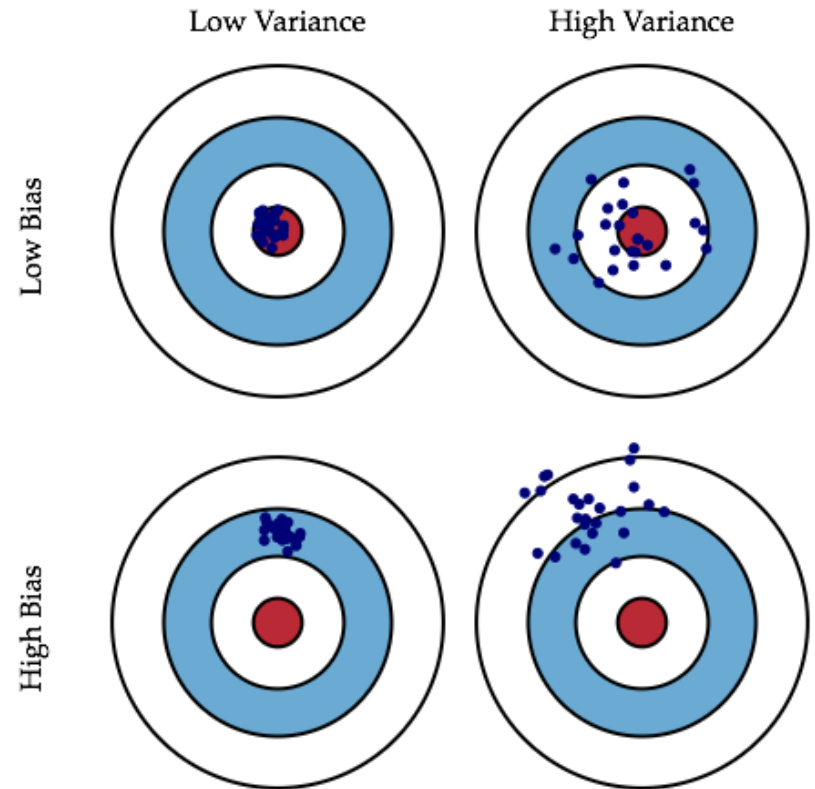
params = [
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001], 'kernel':
 ['rbf']},
]

estimator = GridSearchCV(SVC(), params)
estimator.fit(dataset.data, dataset.target)
```

# Error: Bias vs Variance

**Bias:** the difference between expected (average) prediction of the model and the correct value.

**Variance:** how the predictions for a given point vary between different realizations for the model.



<http://scott.fortmann-roe.com/docs/BiasVariance.html>

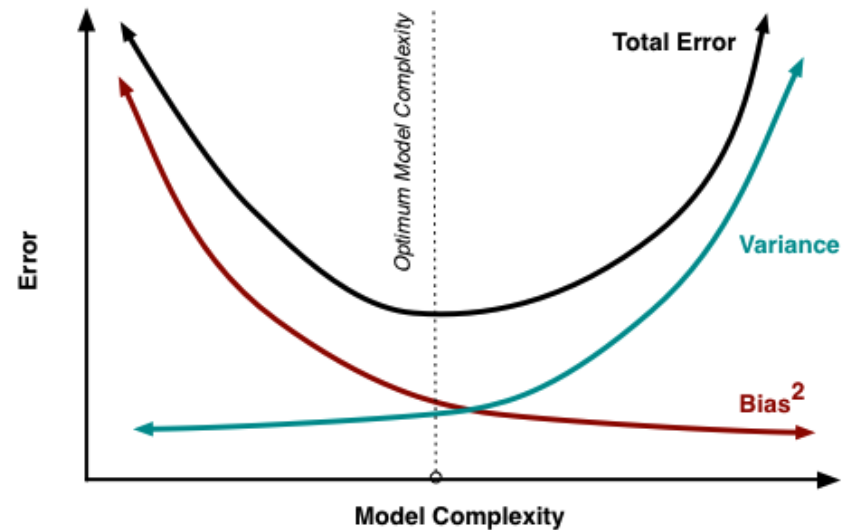
# Bias vs. Variance Trade-Off

Related to model complexity:

The more parameters added to the model (the more complex), Bias is reduced, and variance increased.

## Sources of complexity:

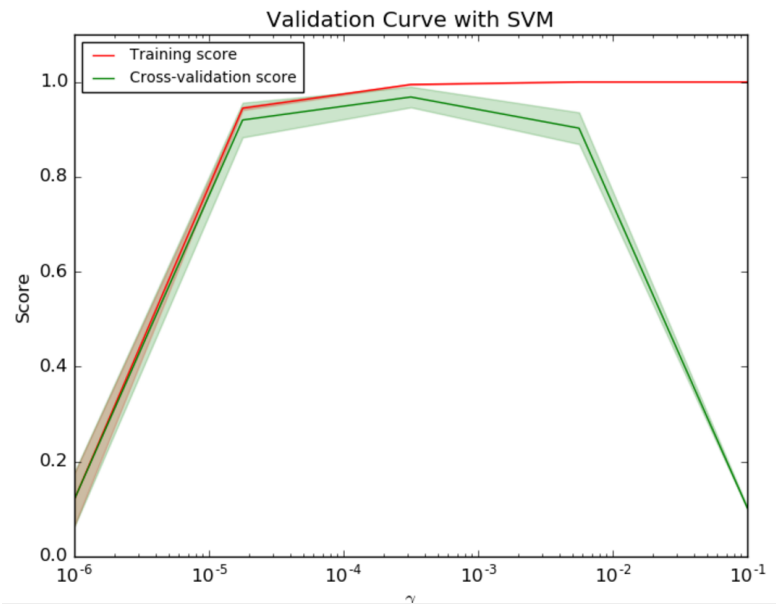
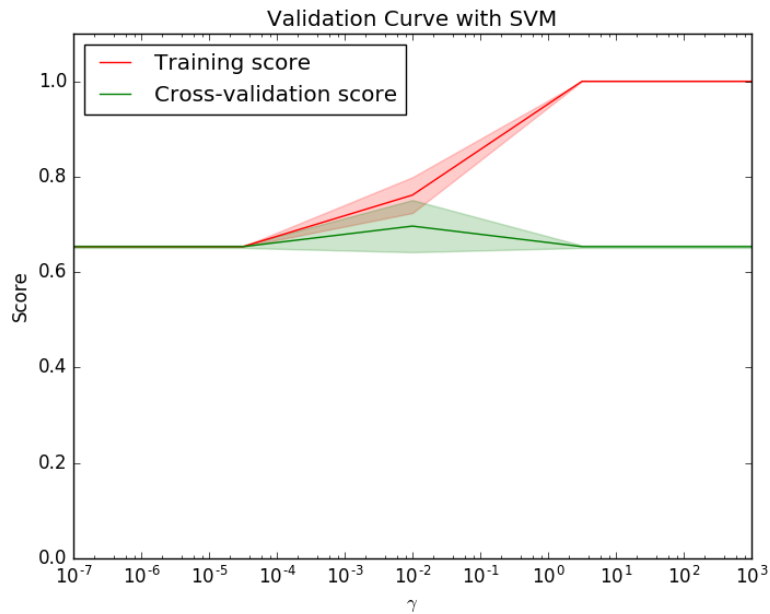
- k (nearest neighbors)
- epochs (neural nets)
- # of features
- learning rate



<http://scott.fortmann-roe.com/docs/BiasVariance.html>



# Visual Parameter Tuning



# Visual Parameter Tuning

- 2 different datasets: tic-tac-toe (left) & digits (right)
- Training vs. validation scores of SVM
- Different values of the kernel parameter gamma
- Things to look for:
  - Training score and validation score both low => Underfit
  - Training score high and validation score low => Overfit

# Lasso

- Reducing bias is one thing, but what if the coefficients are very sparse? E.g. the more dimensions we add, the more *space* goes into the model.
- Lasso prefers fewer parameters attempting to reduce the number of variables the solution depends on.

$$\min_w \frac{1}{2n_{\text{samples}}} \left( \sum (Xw - y)^2 \right) + \alpha \|w\|_1$$

- The term  $\alpha \|w\|_1$  is the L1 norm, whereas in ridge we used the L2 norm,  $\alpha \|w\|_2^2$ .
- See also Least Angle Regression (LARS) as similar.
- Can also use `LassoCV` and `LassoLarsCV`

# Lasso Regression

```
clf = linear_model.Lasso(alpha=0.5)
clf.fit(X_train, y_train)
```

```
Lasso(alpha=0.5, copy_X=True, fit_intercept=True,
      max_iter=1000, normalize=False, positive=False,
      precompute='auto', tol=0.0001, warm_start=False)
```

```
print mean_squared_error(y_test, clf.predict(X_test))
18.84667821
```

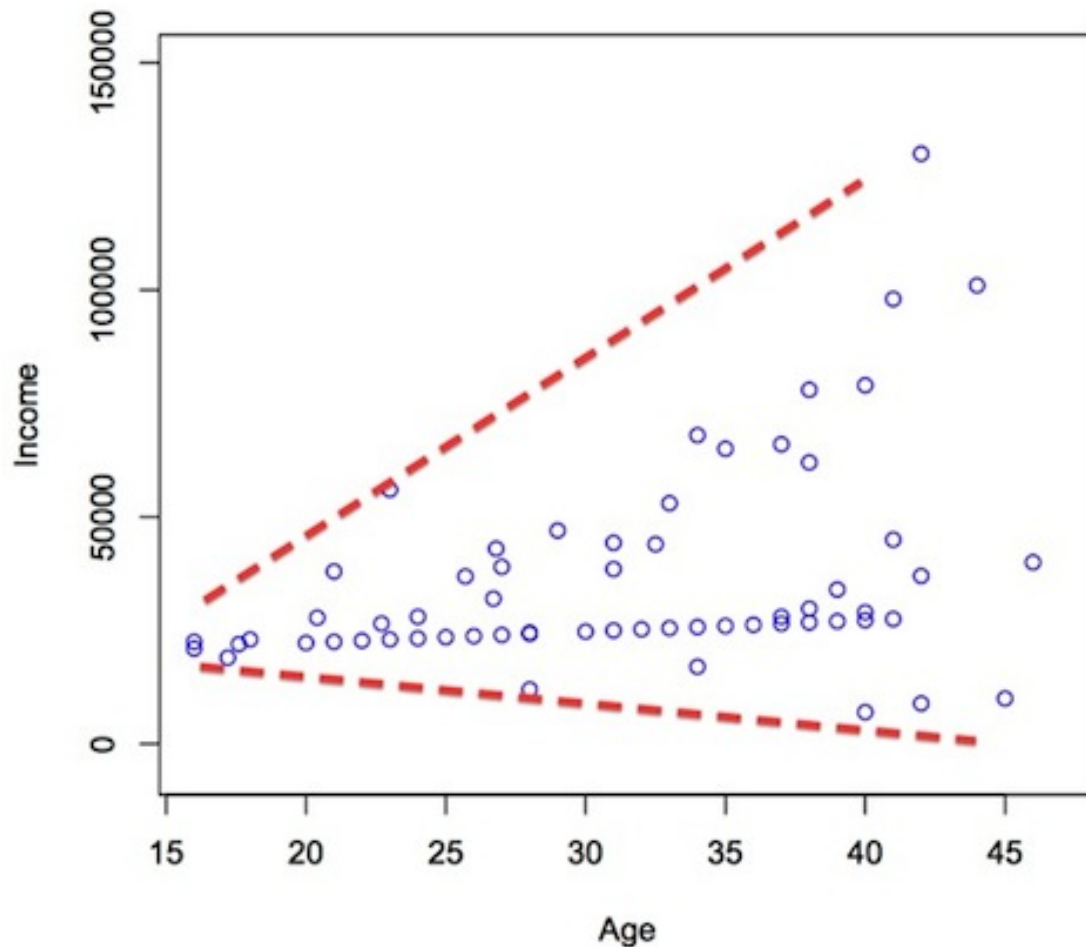
```
clf.score(X_test, y_test)
0.82870491763341947
```

# Instance Variance

**Heteroscedasticity:** variability of variable is unequal along range of predicted values.

**Homoscedasticity:** variance is equal along prediction (assumed in most models).

# Instance Variance



# And More Models

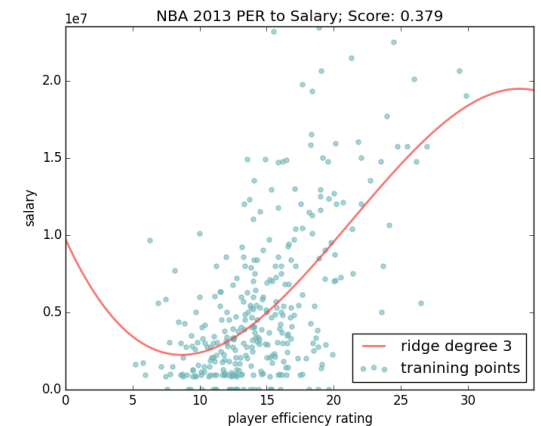
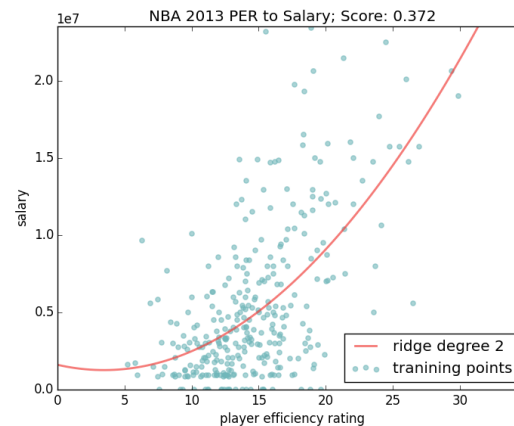
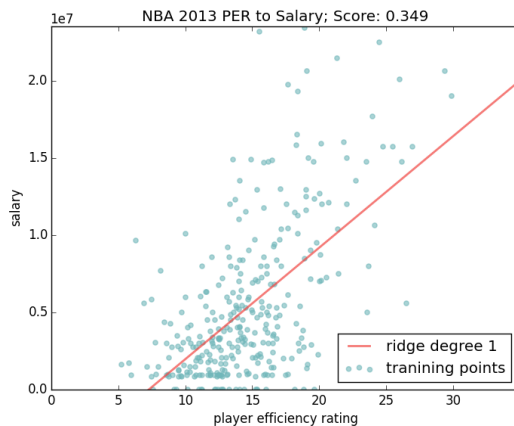
Listed only from the Documentation (not API):

- ElasticNet
- Multi-Task Lasso
- Least Angle Regression
- LARS Lasso
- Orthogonal Matching Pursuit (OMP)
- Bayesian Regression
- Automatic Relevance Determination (ARD)
- Logistic Regression
- Stochastic Gradient Descent
- Perceptron
- Random Sample Consensus (RANSAC)

# Polynomial Regression

In order to do higher order polynomial regression, we can use *linear models* trained on *nonlinear* functions of data!

- Speed of linear model computation
- Fit a wider range of data or functions
- But remember: polynomials aren't the only functions to fit





# Polynomial Regression

The way this works is via *Pipelining*.

Consider the standard linear regression case:

$$\hat{y}(w, x) = w_0 + \sum_{i=1}^n w_i x_i$$

The quadratic case (polynomial degree = 2) is:

$$\hat{y}(w, v, x) = w_0 + \sum_{i=1}^n w_i x_i + \sum_{i=1}^n v_i x_i^2$$

But this can just be seen as a new feature space:

And this feature space can be computed in a linear fashion.

We just need some way to add our 2<sup>nd</sup> degree dimensions.

# Pipelined Model

```
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(PolynomialFeatures(2), linear_model.Ridge())
model.fit(X_train, y_train)
```

```
Pipeline(steps=[('polynomialfeatures',
                  PolynomialFeatures(degree=2, include_bias=True,
                                      interaction_only=False)),
                ('ridge',
                  Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                        max_iter=None, normalize=False, solver='auto',
                        tol=0.001))])
```

```
mean_squared_error(y_test, model.predict(X_test))
3.1498887586451594
```

```
model.score(X_test, y_test)
0.97090576345108104
```

# Pipelines (Steps)

`sklearn.pipeline.Pipeline`

- Sequentially apply *repeatable* transformations to final estimator that can be validated at every step.
- Each step (except for the last) must implement `Transformer`, e.g. `fit` and `transform` methods.
- Pipeline itself implements both methods of `Transformer` and `Estimator` interfaces.

# Transformers

```
class Transformer(Estimator):  
  
    def transform(self, X):  
        """Transforms the input data. """  
        # transform ``X`` to ``X_prime``  
        return X_prime  
  
from sklearn import preprocessing  
  
Xt = preprocessing.normalize(X) # Normalizer  
Xt = preprocessing.scale(X)     # StandardScaler  
  
imputer =Imputer(missing_values='Nan',  
                 strategy='mean')  
Xt = imputer.fit_transform(X)
```

# Scikit-Learn Estimator API

```
class Estimator(object):  
  
    def fit(self, X, y=None):  
        """Fits estimator to data. """  
        # set state of ``self``  
        return self  
  
    def predict(self, X):  
        """Predict response of ``X``. """  
        # compute predictions ``pred``  
        return pred
```

```
from sklearn.pipeline import Pipeline
from sklearn.cross_validation import KFold

pipeline = Pipeline([
    ('extract_essays', EssayExtractor()),
    ('counts', CountVectorizer()),
    ('tf_idf', TfidfTransformer()),
    ('classifier', MultinomialNB())
])

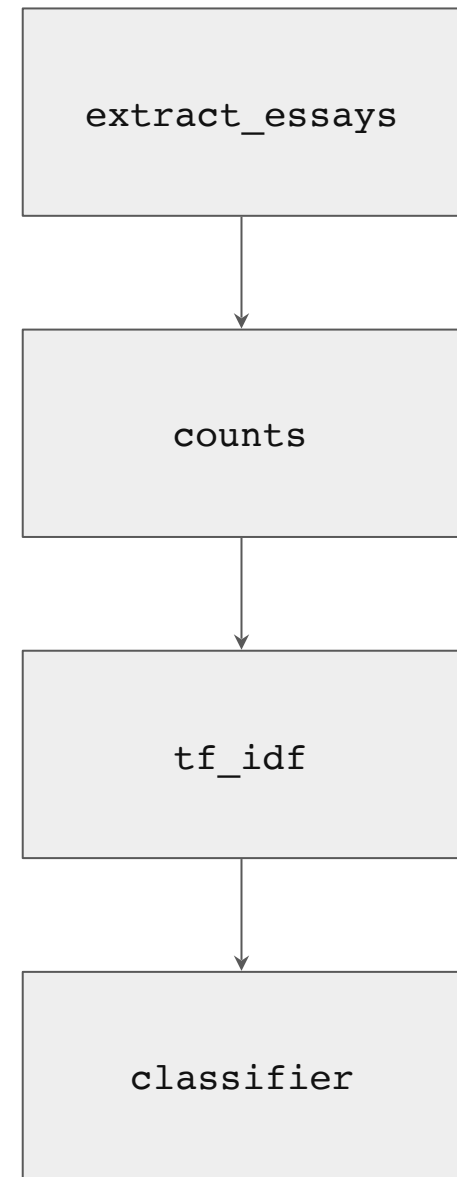
scores = []
folds = KFold(

    n = dataset.data.shape[0], n_folds=12, shuffle=True

)

for tid, cid in folds:
    pipeline.fit(dataset.data[tid], dataset.target[cid])
    score = pipeline.score(dataset.data[cid],
dataset.target[cid])
    scores.append(score)

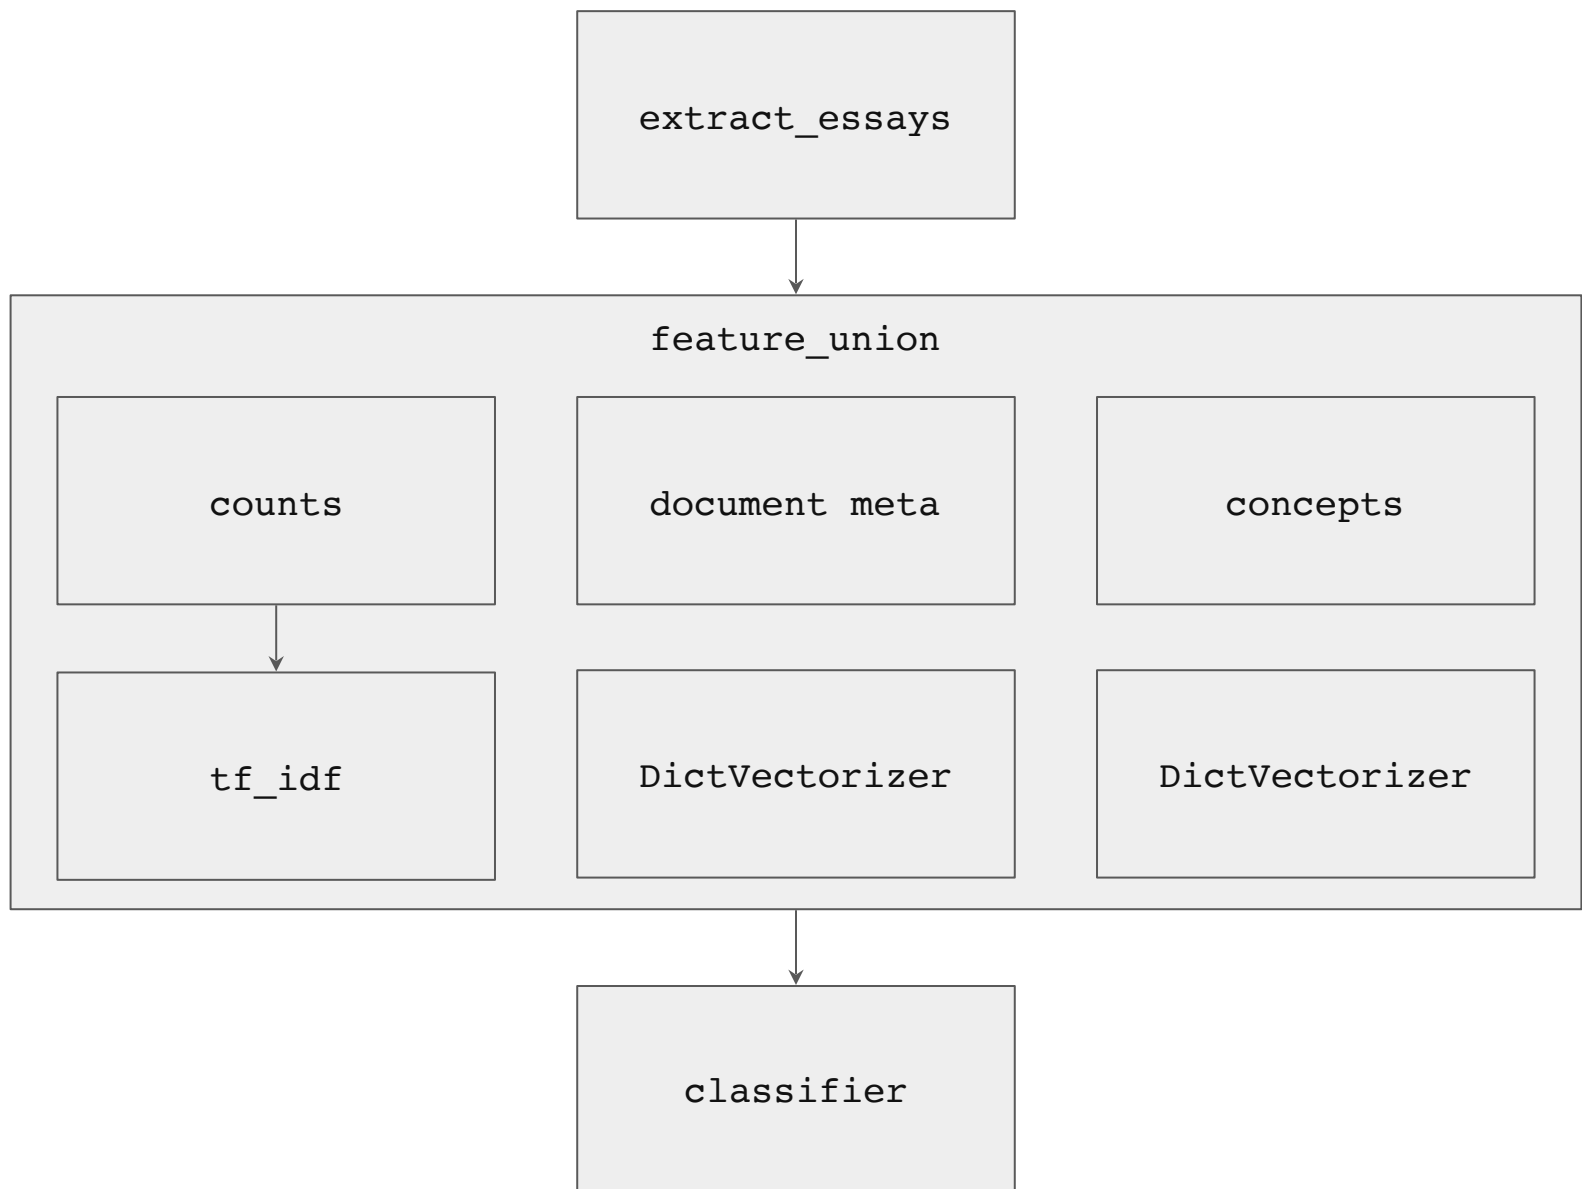
print("Score: {}".format(np.mean(scores)))
```



# Pipelined Feature Extraction

The most common use for the Pipeline is to combine multiple feature extraction methodologies into a single, repeatable processing step.

- FeatureUnion
- SelectKBest
- TruncatedSVD
- DictVectorizer



Feature unions example from Zac's post



# Regression at Scale

# What Makes Scikit-Learn Special

# The Scikit-Learn API

Object-oriented interface centered around the concept of an Estimator:

“An estimator is any object that learns from data; it may be a classification, regression or clustering algorithm or a transformer that extracts/filters useful features from raw data.”

- Scikit-Learn Tutorial

# The Scikit-Learn API

```
class Estimator(object):

    def fit(self, X, y=None):
        """
        Fits estimator to data.
        """
        # set state of self
        return self

    def predict(self, X):
        """
        Predict response of X
        """
        # compute predictions pred
        return pred
```

Buitinck, Lars, et al. "API design for machine learning software: experiences from the scikit-learn project." arXiv preprint arXiv:1309.0238 (2013).

# The Scikit-Learn API

```
class Transformer(Estimator):  
  
    def transform(self, X):  
        """  
        Transforms the input data.  
        """  
        # transform X to X_prime  
        return X_prime
```

Buitinck, Lars, et al. "API design for machine learning software: experiences from the scikit-learn project." arXiv preprint arXiv:1309.0238 (2013).

# The Scikit-Learn API

```
class Pipeline(Transformer):  
  
    @property  
    def named_steps(self):  
        """  
        Returns a sequence of estimators  
        """  
        return self.steps  
  
    @property  
    def _final_estimator(self):  
        """  
        Terminating estimator  
        """  
        return self.steps[-1]
```

Buitinck, Lars, et al. "API design for machine learning software: experiences from the scikit-learn project." arXiv preprint arXiv:1309.0238 (2013).

# The Model Selection Triple



A Venn diagram consisting of three overlapping circles. The leftmost circle is light blue and contains the text 'Feature Analysis'. The middle circle is a medium blue and contains the text 'Algorithm Selection'. The rightmost circle is a dark blue and contains the text 'Hyperparameter Tuning'. The circles overlap in a way that creates a central region where all three overlap, and two other regions where two circles overlap. The text is white and centered within each circle.

Feature Analysis

Algorithm Selection

Hyperparameter Tuning



# The Model Selection Triple



## Feature Analysis

- **Define** a bounded, high dimensional feature space that can be effectively modeled.
- **Transform** and manipulate the space to make modeling easier.
- **Extract** a feature representation of each instance in the space.

# The Model Selection Triple



## Algorithm Selection

- Select a **model family** that best/correctly defines the relationship between the variables of interest.
- Define a **model form** that specifies exactly how features interact to make a prediction.
- Train a **fitted model** by optimizing internal parameters to the data.

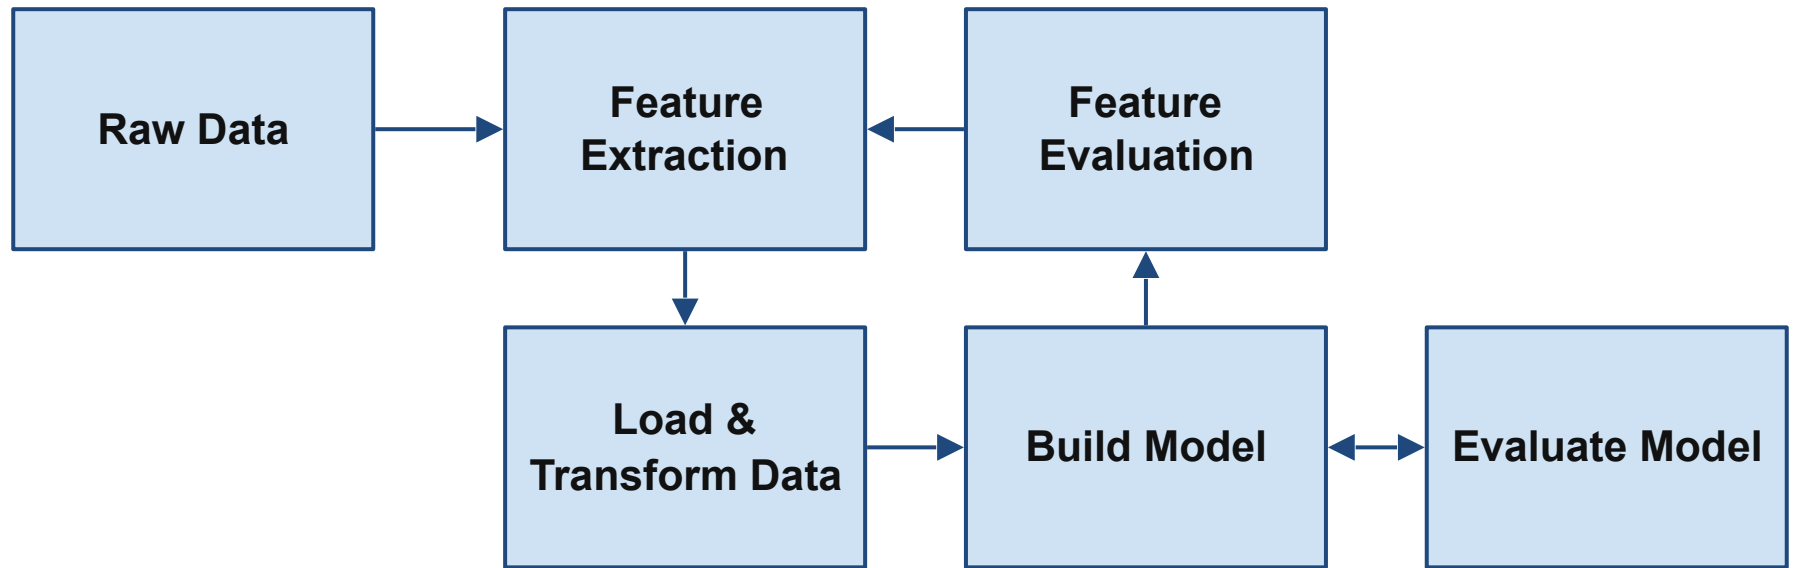
# The Model Selection Triple



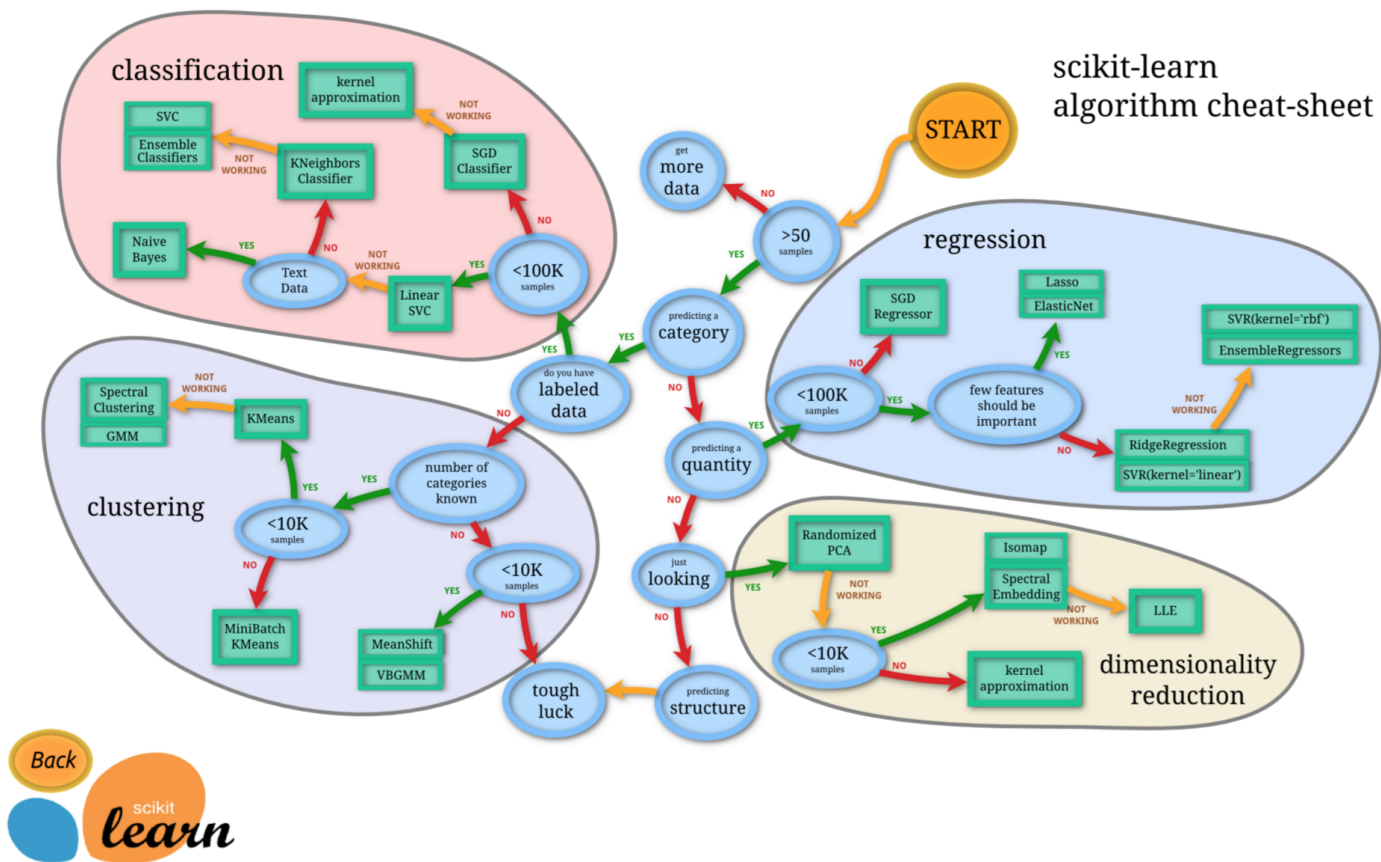
## Hyperparameter Tuning

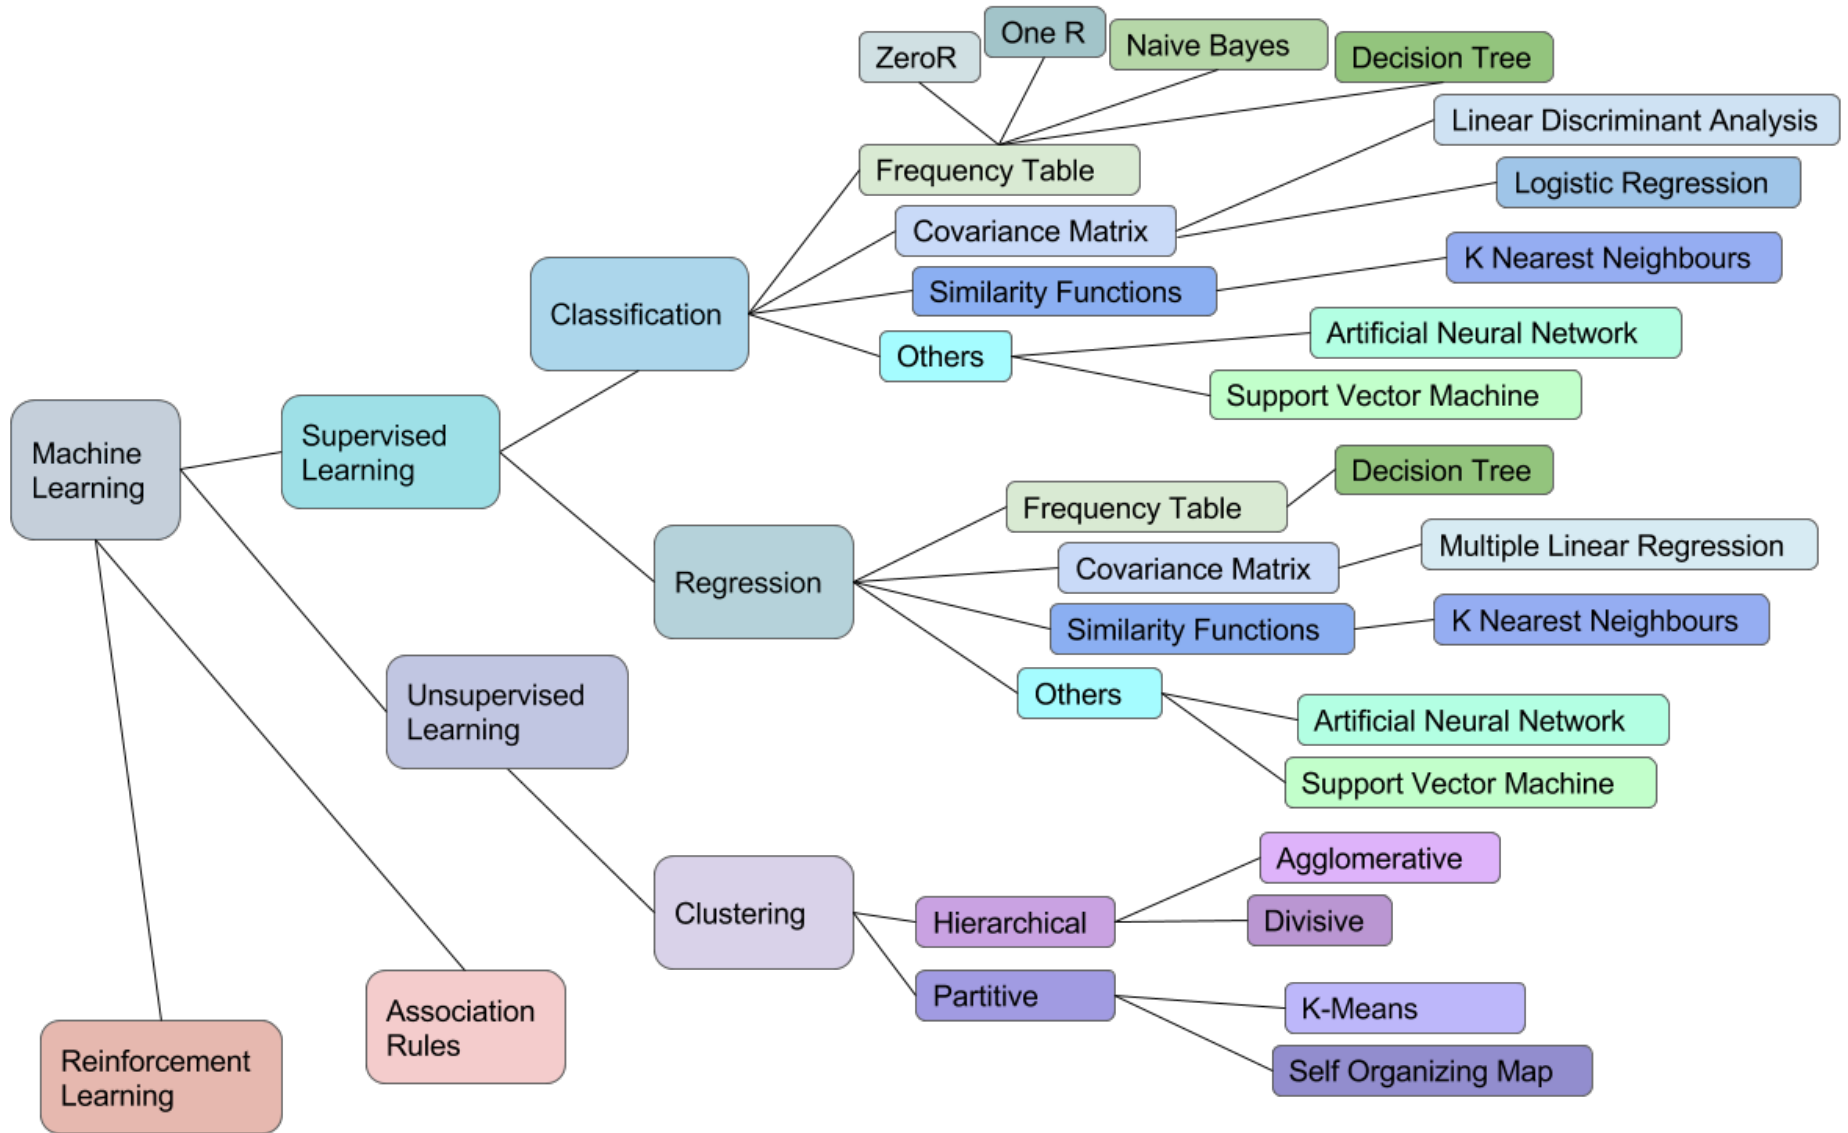
- **Evaluate** how the model form is interacting with the feature space.
- **Identify** hyperparameters (parameters that affect training or the prior, not prediction)
- **Tune** the fitting and prediction process by modifying these params.

# Preliminary Workflow



# Choosing the Right Estimator





# Spark MLlib

# Spark MLlib

Spark's scalable machine learning library consisting of common learning algorithms and utilities, including classification, regression, clustering, collaborative filtering, dimensionality reduction, as well as underlying optimization primitives.



# Spark MLlib: Highlights

- Summary statistics and correlation
- Hypothesis testing, random data generation
- Linear models of regression (SVMs, logistic and linear regression)
- Naive Bayes and Decision Tree classifiers
- Collaborative Filtering with ALS
- K-Means clustering
- SVD (singular value decomposition) and PCA
- Stochastic gradient descent

# Summary of Spark Regression Models

- Linear regression
- Generalized linear regression
- Decision tree regression
- Random forest regression
- Gradient-boosted tree regression
- Survival regression
- Isotonic regression

# Hands-On Lab

# Tasks

- Linear regression
- Generalized linear regression
- Decision tree regression
- Random forest regression
- Gradient-boosted tree regression
- Survival regression
- Isotonic regression